

Technical Overview

Smart i-LAB Testbed

Frontend, backend, architecture, CI/CD, ML training, and operations

`pctiope/smart-i-lab-testbed`

May 2026

Objective

Review question

How does the combined Smart i-LAB repository move from lab devices to API, frontends, model training, deployment, and operational verification?

Repo-grounded

No runtime secrets

No generated model data

- Keep the Metropolis Beamer base, but use a darker presentation rhythm.
- Show flows and ownership boundaries instead of dense path tables.
- Separate IoT1 cutover status from Zone 5 CI/CD and ML delivery.
- Describe Zone 5 and AIR1 model contracts without inventing metrics.



01

Frontend



Tracked Digital Twin floor asset.


Purpose

- Web visualization for Room 308 device status and selected controls.
- Demonstrates the IoT1 REST API as the user-facing integration layer.
- Served as static assets by nginx after a Vite build.

Boundary

- API base is configured at build time with `VITE_API_URL`.
- Three.js assets stay with the tracked frontend source.
- It does not own model inference or data collection.

Zone 5 Dashboard




What the operator sees

- Occupancy probability, CV truth, recent history, squared error, and sensor context.
- Annotated MJPEG from the file-backed latest frame path.
- Split-pane layout with the camera feed beside the readout surface.

Backend contract

`/api/health`, `/api/current`, `/api/history`, `/api/stream`, and `/api/video.mjpg` are preserved so UI changes do not redefine runtime APIs.

AIR1 All-Zones Dashboard



Operational view

- One shared AIR1 model reports per-zone probabilities for zones 1–15.
- Two camera streams feed separate zone masks and one per-zone label table.
- The app can remain up while model readiness is blocked by a missing production pointer.

Important separation

- AIR1 excludes power and mmWave from model inputs.
- Missing camera coverage remains null; it is not converted to vacant.




02

Backend

IoT1 REST API

Role

- Express API for sensor reads, device registries, transactions, groups, logs, and actuator paths.
- The Digital Twin and CV packages consume the API instead of talking directly to device tables.
- Device histories are stored in PostgreSQL with TimescaleDB hypertables.



Security And Storage

Auth

- API-key checks.
- Access-level validation.
- CV clients continue to send x-api-key.

Hardening

- Parameterized SQL or allow-lists.
- Helmet, rate limits, body caps.
- Non-root containers.

Storage

- Versioned migrations.
- Sensor tables created lazily.
- Dirty legacy tables remain cleanup work.

`docs/IOT1_AUDIT.md`

`docs/IOT1_REMEDIATION.md`


`migrations/apply.py`




03

Architecture

User And Data Flow



End-To-End Architecture



Device data

REST integration

CV labels


Model delivery



04

CI/CD

GitHub Workflow Topology



Workflow

Ownership

`zone5-systemd-source-deploy.yml`

Production source deploy; restarts live services only.

`zone5-compose-ci-cd.yml`

Staging Compose path; self-hosted deploy to separate ports.

`zone5-production-retrain.yml`

Starts the trainer and packages a promoted model only after gating.

`zone5-model-delivery.yml`

Installs a vetted model artifact and updates the production pointer.

Staging Compose Vs. Production systemd

Production

- Branch: `main`
- Runtime: user-level `systemd`
- Ports: backend 8000, Vite 8015
- Source deploy refuses active trainer collisions.

Staging

- Branch: `cicd/zone5-compose`
- Runtime: Docker Compose
- Ports: backend 8005, Vite 8016
- Copies env and model pointer from production when needed.

Source deploy, staging deploy, training, and model delivery stay separate.

Retrain And Model Delivery



- Scheduled retraining is present but disabled until the manual proof run is accepted.
- Promotion failure leaves collection and serving alone.
- Sanitized model registry metadata can be committed; generated run directories stay off-source.



05

ML Training

CV Labels

Zone 5

- One camera, one desk mask, one 10-second Zone 5 label stream.
- Target is `zone_occupied`.
- Occupied when median person count crosses the configured threshold.

AIR1 all-zones

- Two cameras, per-zone masks, long-form labels keyed by `timestamp + zone_id`.
- Target is `occupied`.
- Missing camera coverage remains null instead of becoming a negative label.



Zone 5 Methodology


Feature families

- AIR-1 Zone 5: temperature, humidity, CO2, PM2.5.
- Smart plug: `power_s5`.
- mmWave: `mmwave_s5` and recency context.
- SEN55: particulate, environment, VOC, NO_x.
- Time: hour and day-of-week cycles.

Training contract


- Current marker:
`zone5_mmwave_recency_v1`.
- Missing indicators are diagnostics and metadata, not model inputs.
- Legacy
`zone5_missingness_decoupled_v1`
artifacts remain loadable for live compatibility.
- Trainer snapshots the rolling CSV before fitting.

mmWave Recency Contract



- mmWave remains an input feature, not the training target.
- Recency features are computed from prior samples only.
- The model sees short-history motion context without future leakage.
- Dashboard history surfaces the lookback context beside probability.

AIR1 All-Zones Pipeline



Contract guard

zone_id is retained but not an input

Power and mmWave are excluded

Production pointer gates inference

Why it matters

It supports grouping, splits, audit, and display without becoming a shortcut feature.

AIR1 parity does not import Zone 5-specific signals.


The UI can be healthy before a promoted model exists.



06

Deployment

Deployment Architecture



Runtime credentials, generated data, model runs, and logs live on the server, not in the presentation or source history.

Verification And Performance Gates

CI gates

- Python compile checks for package and app modules.
- Zone 5 ground-truth contract tests.
- Vite build checks.
- Systemd unit verification.
- Docker Compose config and image build.

Runtime gates

- `/api/health` confirms app health and model readiness fields.
- Model promotion checks smoke loading and non-regression conditions.
- IoT1 cutover still requires backup, migrations, deploy, smoke test, and log tail.
- Latency and label-balance monitoring remain roadmap items.

Roadmap And Risks

Near-term roadmap

1. Complete the IoT1 production cutover sequence.
2. Dedupe legacy per-device timestamp histories.
3. Prove manual Zone 5 retraining before enabling scheduled retrain.
4. Add model freshness, promotion, and upstream API latency monitoring.

Risks to keep visible

- Live QA depends on lab network, TimescaleDB, EMQX, Home Assistant, cameras, and devices.
- Training quality depends on both-class CV labels across usable calendar windows.
- Contract changes require artifact compatibility and dashboard context updates.

Source Anchors

Topic	Primary source paths
IoT1 status and cutover	docs/STATUS.md, docs/IOT1_PRE_DEPLOYMENT.md, docs/IOT1_REMEDIATION.md
Digital Twin	Smart-iLab_DigitalTwin/, README.md
Zone 5 ML and operations	zone5_cv_time_features_package/PIPELINE.md, zone5/.../feature_contract.py
AIR1 all-zones ML and operations	air1_all_zones_cv_time_features_package/PIPELINE.md, air1_all_zones/.../feature_contract.py
CI/CD and deployment	.github/workflows/, zone5/.../CICD_SYSTEMD.md, zone5/.../CICD_DOCKER_COMPOSE.md

Closing point: source, staging runtime, production runtime, and model artifacts each have a separate owner and gate.